

Movie_Similarity_Model

December 14, 2016

```
In [1]: import heapq
import math
import time
from collections import defaultdict
from collections import namedtuple
from contextlib import contextmanager
from datetime import datetime

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split

%matplotlib inline
```

```
In [5]: @contextmanager
def elapsed_time(title):
    start = time.time()
    yield
    elapsed = time.time() - start
    print '%s: %.2f secs' % (title, elapsed)

def get_xy(ratings_df):
    y = ratings_df['rating']
    x = ratings_df.drop('rating', axis=1)
    return x, y

def root_mean_squared_error(y, y_pred):
    return math.sqrt(mean_squared_error(y, y_pred))

def show_scores_plot(k_neighbors_values, val_scores, train_scores, model_na
_, ax = plt.subplots(1, 1, figsize=(15, 10))
```

```

ax.plot(k_neighbors_values, val_scores, label='validation')
ax.plot(k_neighbors_values, train_scores, label='train')

ax.set_xlabel('k_neighbors')
ax.set_ylabel('$R^2$')
ax.set_title('Test and validation scores for different k_neighbors values')

ax.legend(loc='best')

plt.tight_layout()
plt.show()

```

```

def score_model(ratings_df, model_f, model_name):
    train_val_ratings_df, test_ratings_df = train_test_split(ratings_df)

    train_ratings_df, validation_ratings_df = train_test_split(train_val_ratings_df)

    best_score = -float('inf')
    best_k_neighbors = None

    model = model_f()

    model = model.fit(train_ratings_df)

    k_neighbors_values = [1, 5, 10, 20, 30, 40, 50, 60, 80, 100]

    val_scores = []
    train_scores = []

    for k_neighbors in k_neighbors_values:
        model.set_k_neighbors(k_neighbors=k_neighbors)

        x_train, y_train = get_xy(train_ratings_df)
        x_val, y_val = get_xy(validation_ratings_df)

        y_train_pred = model.predict(x_train)
        y_val_pred = model.predict(x_val)

        train_score = r2_score(y_train, y_train_pred)
        val_score = r2_score(y_val, y_val_pred)

        if val_score > best_score:
            best_score = val_score
            best_k_neighbors = k_neighbors

    val_scores.append(val_score)

```

```

train_scores.append(train_score)

    print 'k: %d, validation score: %.5f, train score: %.5f\n' % (k_neighbors, validation_score, train_score)

print 'best k: %d, best score: %.5f' % (best_k_neighbors, best_score)

model = model_f(k_neighbors=best_k_neighbors)

model = model.fit(train_val_ratings_df)

x_train_val, y_train_val = get_xy(train_val_ratings_df)
x_test, y_test = get_xy(test_ratings_df)

y_train_val_pred = model.predict(x_train_val)
y_test_pred = model.predict(x_test)

train_val_score = r2_score(y_train_val, y_train_val_pred)
test_score = r2_score(y_test, y_test_pred)

train_val_rmse = root_mean_squared_error(y_train_val, y_train_val_pred)
test_rmse = root_mean_squared_error(y_test, y_test_pred)

print 'train score: %.4f, test score: %.4f' % (train_val_score, test_score)
print 'train rmse: %.4f, test rmse: %.4f' % (train_val_rmse, test_rmse)

show_scores_plot(k_neighbors_values, val_scores, train_scores, model_name)

```

```

In [6]: def date_parse(time_in_secs):
        return datetime.utcfromtimestamp(float(time_in_secs))

def read_ratings_df_with_timestamp(file_name):
    with elapsed_time('loaded csv'):
        ratings_df = pd.read_csv(file_name, parse_dates=['timestamp'], date_format='%Y-%m-%d %H:%M:%S')
    return ratings_df

In [7]: class BaselineModel(object):
        def predict_rating(self, user_id, movie_id):
            pass

        def predict(self, x):
            return [self.predict_rating(row['userId'], row['movieId']) for _, row in x.iterrows()]

        def score(self, x, y):
            return r2_score(y, self.predict(x))

class BaselineEffectsModel(BaselineModel):

```

```

def __init__(self, movie_lambda=5.0, user_lambda=20.0):
    self.movie_lambda = movie_lambda
    self.user_lambda = user_lambda

    self.y_mean = None
    self.movie_effects = None
    self.user_effects = None
    self.user_groups = None

def calculate_movie_effect(self, ratings):
    return (ratings - self.y_mean).sum() / (self.movie_lambda + len(ratings))

def calculate_movie_effects(self, movie_ratings):
    return movie_ratings.agg(lambda ratings: self.calculate_movie_effect(ratings))

def calculate_user_effect(self, ratings_df):
    s = 0.0
    for _, row in ratings_df.iterrows():
        s += row['rating'] - self.y_mean - self.movie_effects[row['movieId']]

    return s / (self.user_lambda + len(ratings_df))

def calculate_user_effects(self, user_groups):
    user_ids = []
    user_effects = []

    for user_id, group in user_groups:
        user_effect = self.calculate_user_effect(group)

        user_ids.append(user_id)
        user_effects.append(user_effect)

    return pd.Series(user_effects, index=user_ids)

def fit(self, ratings_df):
    with elapsed_time('effects init'):
        _, y_train = get_xy(ratings_df)
        self.y_mean = y_train.mean()

        movie_ratings = ratings_df.groupby('movieId')['rating']
        self.user_groups = ratings_df.groupby('userId')

        self.movie_effects = self.calculate_movie_effects(movie_ratings)
        self.user_effects = self.calculate_user_effects(self.user_groups)

    return self

def create_modified_ratings(self, ratings_df):

```

```

ratings_df = ratings_df.copy()

for index, row in ratings_df.iterrows():
    user_id = row['userId']
    movie_id = row['movieId']
    rating = row['rating']
    pred_rating = self.predict_baseline_rating(user_id, movie_id)

    residual = rating - pred_rating

    ratings_df.loc[index, 'rating'] = residual

return ratings_df

def predict_baseline_rating(self, user_id, movie_id):
    return self.y_mean + self.movie_effects.get(movie_id, 0.0) + self.u

def predict_rating(self, user_id, movie_id):
    return self.predict_baseline_rating(user_id, movie_id)

```

Movie similarity model.

First we removed all main global effects the same way we did for our baseline model. We subtracted the total rating mean, then we removed the movie effects and then user effects.

As a result our utility matrix was the residuals after applying the baseline model.

It allowed us to remove some scale differences in the way different users rate movies.

In our movie similarity model we predict the rating for a movie (movie_id) and a user (user_id) by

- 1) finding k closest neighbors (k=40,50)

For assessing item-item similarity we used a distance based on the mean squared error between items [1]:

$$s_{ij} = \frac{|U(i,j)|}{\sum_{u \in U(i,j)} (r_{ui} - r_{uj})^2 + \alpha}, \text{ where } U(i,j) \text{ is the set of users who rated both items } j \text{ and } k.$$

- 2) since relatively large number of movies have low number of ratings (1-3), some of the movies have zero neighbors, in this case we use a baseline prediction (in 6.9-7.2% of the cases for the test set).

[1] R.Bell, Y.Koren, "Improved Neighborhood-based Collaborative Filtering", *KDD-Cup and Workshop*, ACM press, 2007

```
In [8]: MovieSimilarity = namedtuple('MovieSimilarity', ['movie_id', 'similarity'])
```

```

class MovieSimilarityModel(BaselineModel):
    def __init__(self, k_neighbors=40):
        self.k_neighbors = k_neighbors

        self.baseline_model = BaselineEffectsModel()

```

```

self.ratings_by_movie = defaultdict(dict)
self.ratings_by_user = defaultdict(dict)
self.raters_by_movie = {}
self.movie_similarity = {}
# self.movie_aj = {}

def set_k_neighbors(self, k_neighbors):
    self.k_neighbors = k_neighbors

def calculate_common_raters(self, movie_id_1, movie_id_2):
    raters1 = self.raters_by_movie[movie_id_1]
    raters2 = self.raters_by_movie[movie_id_2]
    return raters1 & raters2

def get_common_ratings(self, movie_id, raters):
    all_ratings = self.ratings_by_movie[movie_id]
    ratings = []
    for rater_id in raters:
        ratings.append(all_ratings[rater_id])

    return np.array(ratings)

def calculate_similarity(self, movie_id_1, movie_id_2):
    common_raters = self.calculate_common_raters(movie_id_1, movie_id_2)
    support = len(common_raters)
    if support <= 1:
        similarity = 0.0
        # aij = 0.0
    else:
        ratings1 = self.get_common_ratings(movie_id_1, common_raters)
        ratings2 = self.get_common_ratings(movie_id_2, common_raters)

        alpha = 4.0

        similarity = support / (np.power(ratings1 - ratings2, 2).sum())

        # aij = np.multiply(ratings1, ratings2).sum() / support

    return similarity

def fit(self, ratings_df):
    with elapsed_time('fit'):
        self.baseline_model.fit(ratings_df)

    ratings_df = self.baseline_model.create_modified_ratings(ratings_df)

    unique_movie_ids = np.array(sorted(ratings_df['movieId']).unique())

```

```

        for _, row in ratings_df.iterrows():
            movie_id = row['movieId']
            user_id = row['userId']
            rating = row['rating']
            self.ratings_by_movie[movie_id][user_id] = rating
            self.ratings_by_user[user_id][movie_id] = rating

        for movie_id in unique_movie_ids:
            self.raters_by_movie[movie_id] = set(self.ratings_by_movie[movie_id].keys())

        for movie_index_1, movie_id_1 in enumerate(unique_movie_ids):
            for movie_index_2 in xrange(movie_index_1 + 1, len(unique_movie_ids)):
                movie_id_2 = unique_movie_ids[movie_index_2]

                similarity = self.calculate_similarity(movie_id_1, movie_id_2)
                movie_pair = (movie_id_1, movie_id_2)
                self.movie_similarity[movie_pair] = similarity
                # self.movie_aij[movie_pair] = aij

        return self

    def get_similarity(self, movie_id_1, movie_id_2):
        if movie_id_1 < movie_id_2:
            id_1 = movie_id_1
            id_2 = movie_id_2
        else:
            id_1 = movie_id_2
            id_2 = movie_id_1

        return self.movie_similarity.get((id_1, id_2), -1.0)

    def clear_predict_caches(self):
        self.zero_prediction_count = 0

    def predict_rating(self, user_id, movie_id):
        ratings = self.ratings_by_user[user_id]

        elements = []

        for movie_id_2 in ratings:
            if movie_id != movie_id_2:
                similarity = self.get_similarity(movie_id, movie_id_2)
                if similarity > 0.0:
                    elements.append(MovieSimilarity(movie_id_2, similarity))

        movie_similarities = heapq.nlargest(self.k_neighbors, elements, key=lambda x: x.similarity)

        if len(movie_similarities) > 0:

```

```

        similarity_sum = 0.0
        product_sum = 0.0
        for movie_similarity in movie_similarities:
            movie_id_2 = movie_similarity.movie_id
            rating = ratings[movie_id_2]
            similarity = movie_similarity.similarity

            product_sum += similarity * rating
            similarity_sum += similarity

        rating = product_sum / similarity_sum
    else:
        rating = 0.0
        self.zero_prediction_count += 1

    result = self.baseline_model.predict_baseline_rating(user_id, movie_id)

    return result

def predict(self, x):
    self.clear_predict_caches()
    predictions = [self.predict_rating(row['userId'], row['movieId']) for row in x.iterrows()]
    print 'used baseline predictions: %.1f%%' % (100.0 * self.zero_prediction_count / len(predictions))
    return predictions

```

```

ratings_df = read_ratings_df_with_timestamp('ml-latest-small/ratings.csv')

with elapsed_time('score model'):
    score_model(ratings_df, model_f=MovieSimilarityModel, model_name='movie_similarity_model')

```

```

loaded csv: 0.13 secs
effects init: 5.18 secs
fit: 81.15 secs
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 1, validation score: -0.16982, train score: 0.65154

```

```

used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 5, validation score: 0.22471, train score: 0.75729

```

```

used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 10, validation score: 0.26815, train score: 0.73911


```

```

used baseline predictions: 4.9%
used baseline predictions: 9.1%

```


k: 20, validation score: 0.28365, train score: 0.70066
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 30, validation score: 0.28770, train score: 0.67197
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 40, validation score: 0.28723, train score: 0.65002
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 50, validation score: 0.28714, train score: 0.63265
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 60, validation score: 0.28707, train score: 0.61805
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 80, validation score: 0.28558, train score: 0.59584
used baseline predictions: 4.9%
used baseline predictions: 9.1%
k: 100, validation score: 0.28510, train score: 0.57967
best k: 30, best score: 0.28770
effects init: 6.42 secs
fit: 118.84 secs
used baseline predictions: 3.9%
used baseline predictions: 7.0%
train score: 0.6879, test score: 0.3045
train rmse: 0.5921, test rmse: 0.8776



score model: 663.38 secs

```
In [9]: UserSimilarity = namedtuple('UserSimilarity', ['user_id', 'similarity'])
```

```
class UserSimilarityModel(BaselineModel):
    def __init__(self, k_neighbors=40):
        self.k_neighbors = k_neighbors

        self.baseline_model = BaselineEffectsModel()
        self.ratings_by_movie = defaultdict(dict)
        self.ratings_by_user = defaultdict(dict)
        self.movies_by_user = {}
        self.user_similarity = {}

    def set_k_neighbors(self, k_neighbors):
        self.k_neighbors = k_neighbors

    def calculate_common_movies(self, user_id_1, user_id_2):
        movies1 = self.movies_by_user[user_id_1]
        movies2 = self.movies_by_user[user_id_2]
        return movies1 & movies2
```

```

def get_common_ratings(self, user_id, movies):
    all_ratings = self.ratings_by_user[user_id]
    ratings = []
    for movie_id in movies:
        ratings.append(all_ratings[movie_id])

    return np.array(ratings)

def calculate_similarity(self, user_id_1, user_id_2):
    common_movies = self.calculate_common_movies(user_id_1, user_id_2)
    support = len(common_movies)
    if support <= 1:
        similarity = 0.0
        # aij = 0.0
    else:
        ratings1 = self.get_common_ratings(user_id_1, common_movies)
        ratings2 = self.get_common_ratings(user_id_2, common_movies)

        alpha = 4.0

        similarity = support / (np.power(ratings1 - ratings2, 2).sum())

    return similarity

def fit(self, ratings_df):
    with elapsed_time('fit'):
        self.baseline_model.fit(ratings_df)

    ratings_df = self.baseline_model.create_modified_ratings(ratings_df)

    unique_user_ids = np.array(sorted(ratings_df['userId']).unique())

    for _, row in ratings_df.iterrows():
        movie_id = row['movieId']
        user_id = row['userId']
        rating = row['rating']
        self.ratings_by_movie[movie_id][user_id] = rating
        self.ratings_by_user[user_id][movie_id] = rating

    for user_id in unique_user_ids:
        self.movies_by_user[user_id] = set(self.ratings_by_user[user_id].keys())

    for user_index_1, user_id_1 in enumerate(unique_user_ids):
        for user_index_2 in xrange(user_index_1 + 1, len(unique_user_ids)):
            user_id_2 = unique_user_ids[user_index_2]

            similarity = self.calculate_similarity(user_id_1, user_id_2)
            user_pair = (user_id_1, user_id_2)

```

```

        self.user_similarity[user_pair] = similarity

    return self

def get_similarity(self, user_id_1, user_id_2):
    if user_id_1 < user_id_2:
        id_1 = user_id_1
        id_2 = user_id_2
    else:
        id_1 = user_id_2
        id_2 = user_id_1

    return self.user_similarity.get((id_1, id_2), -1.0)

def clear_predict_caches(self):
    self.zero_prediction_count = 0

def predict_rating(self, user_id, movie_id):
    ratings = self.ratings_by_movie[movie_id]

    elements = []

    for user_id_2 in ratings:
        if user_id != user_id_2:
            similarity = self.get_similarity(user_id, user_id_2)
            if similarity > 0.0:
                elements.append(UserSimilarity(user_id_2, similarity))

    user_similarities = heapq.nlargest(self.k_neighbors, elements, key=

    if len(user_similarities) > 0:
        similarity_sum = 0.0
        product_sum = 0.0
        for user_similarity in user_similarities:
            user_id_2 = user_similarity.user_id
            rating = ratings[user_id_2]
            similarity = user_similarity.similarity

            product_sum += similarity * rating
            similarity_sum += similarity

        rating = product_sum / similarity_sum
    else:
        rating = 0.0
        self.zero_prediction_count += 1

    result = self.baseline_model.predict_baseline_rating(user_id, movie

```

```

        return result

    def predict(self, x):
        self.clear_predict_caches()
        predictions = [self.predict_rating(row['userId'], row['movieId']) for row in x.iterrows()]
        print 'used baseline predictions: %.1f%%' % (100.0 * self.zero_predictions / len(predictions))
        return predictions

    with elapsed_time('build model'):
        score_model(ratings_df, model_f=UserSimilarityModel, model_name='user similarity')

effects init: 5.13 secs
fit: 43.29 secs
used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 1, validation score: -0.07280, train score: 0.32184

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 5, validation score: 0.22003, train score: 0.44494

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 10, validation score: 0.24284, train score: 0.42677

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 20, validation score: 0.25189, train score: 0.39777

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 30, validation score: 0.25155, train score: 0.38140

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 40, validation score: 0.25084, train score: 0.37151

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 50, validation score: 0.24993, train score: 0.36488

used baseline predictions: 4.9%
used baseline predictions: 5.0%
k: 60, validation score: 0.24924, train score: 0.36026

used baseline predictions: 4.9%
used baseline predictions: 5.0%

```

k: 80, validation score: 0.24859, train score: 0.35467

used baseline predictions: 4.9%

used baseline predictions: 5.0%

k: 100, validation score: 0.24828, train score: 0.35182

best k: 20, best score: 0.25189

effects init: 6.40 secs


fit: 62.81 secs

used baseline predictions: 3.9%

used baseline predictions: 3.6%

train score: 0.4023, test score: 0.2708

train rmse: 0.8189, test rmse: 0.9006



build model: 236.93 secs

In []: